

# Bayesian Modeling with Stan: An Introductory Workshop

Jakob Torgander

Department of Statistics, Uppsala University  
*jakob.torgander@statistik.uu.se*

November 5, 2025

# Outline

- ① Bayesian modeling
- ② Introduction to Stan
- ③ Model diagnostics and evaluation

## About me/us

- Jakob Torgander, PhD student in Statistics at Uppsala University
- Research in Bayesian statistics & Probabilistic programming
- Part of Interdisciplinary Bayes (InterBayes) network at UU.  
Link: <https://interbayes.github.io>
- Co-developer of posteriordb: a framework for testing and benchmarking Bayesian Inference Algorithms.  
Link: <https://github.com/stan-dev/posteriordb>



- 1 Bayesian modeling
- 2 Introduction to Stan
- 3 Model diagnostics and evaluation

## Bayesian modeling

- Flexible framework for defining and fitting statistical models
- Incorporating prior information → Less data needed
- Framework for updating beliefs/inference as new data is collected
- Intuitive inference. Compare:
  - A Confidence interval:  $\mathbb{P}(\underline{T}(X) < \theta < \bar{T}(X))$ , T is a statistic of data X
  - B Credible interval:  $\mathbb{P}(a < \theta < b)$
- Common bottleneck: computation. Topic of this workshop

## Posterior distribution

Main component of Bayesian modeling: posterior distribution:

$$p(\theta|x) \propto \prod_{i=1}^N \underbrace{f(x_i|\theta)}_{\text{Likelihood}} \cdot \underbrace{p(\theta)}_{\text{Prior}}$$

- $f(x|\theta)$ : relationship between parameter  $\theta$  and data sample  $x$
- $p(\theta)$ : prior beliefs about  $\theta$  (if any)
- Today's task: compute  $p(\theta|x)$  using probabilistic programming

## Why posterior distribution?

Key concept: Posterior contains all information about  $\theta$  given data. Information extracted by computing expectations:

$$\mathbb{E}[h(\theta)] = \int h(\theta)p(\theta|x)d\theta$$

Examples:

$h(\theta)$	$\mathbb{E}[h(\theta)]$
$\theta$	Expected value of $\theta$
$\mathbb{1}(\theta > \tau)$	Probability that $\theta$ is larger than $\tau$ "
$p(y \theta)$	Predictive distribution for a new data point $y$ "

## Example: Normal-normal model

In some (simple) cases, the posterior distribution can be computed analytically. For instance, if we assume the Normal-Normal model

$$\begin{aligned}x|\mu, \sigma^2 &\sim \mathcal{N}(\mu, \sigma^2) && \text{(likelihood)} \\ \mu &\sim \mathcal{N}(\mu_0, \tau_0^2) && \text{(prior)}\end{aligned}$$

Then,

$$\mu|x \sim \mathcal{N}(\mu_{post}, \tau_{post}^2),$$

where

$$\mu_{post} = \frac{\frac{1}{\tau_0^2}\mu_0 + \frac{n}{\sigma^2}\bar{x}}{\frac{1}{\tau_0^2} + \frac{n}{\sigma^2}}, \quad \tau_{post}^2 = \left( \frac{1}{\tau_0^2} + \frac{n}{\sigma^2} \right)^{-1}.$$

## Monte Carlo methods

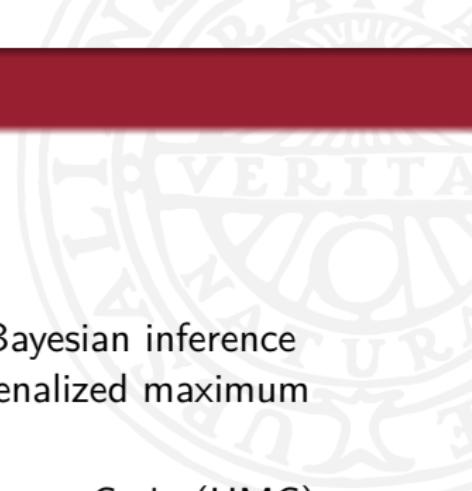
- Problem: posteriors rarely given in closed analytical form.
- Solution: instead of computing full posterior (hard), focus on the expectations (easier)
- Monte Carlo integration: draw samples  $\theta^{(1)}, \dots, \theta^{(N)}$  from the target posterior and approximate:

$$\mathbb{E}[h(\theta)] \approx \frac{1}{N} \sum_{i=1}^N h(\theta^{(i)})$$

- As the number of samples grow  $\frac{1}{N} \sum_{i=1}^N h(\theta^{(i)}) \xrightarrow{a.s} \mathbb{E}[h(\theta)]$
- For drawing the samples  $\theta^{(i)}$ , we will use Stan



- ① Bayesian modeling
- ② Introduction to Stan
- ③ Model diagnostics and evaluation



- Released in 2012
- Probabilistic programming language for Bayesian inference with MCMC, Variational inference and penalized maximum likelihood estimation with optimization
- Implements MCMC using Hamiltonian Monte Carlo (HMC) and the No-U-Turn Sampler (NUTS)
- Written in C++ using similar syntax. Can be used together with R, Python or Julia. (Lab will be in R)
- Link to documentation & tutorials:  
<https://mc-stan.org/docs/>

## Components of Stan program

Recall, we want to compute

$$p(\theta|x) \propto \prod_{i=1}^N \underbrace{f(x_i|\theta)}_{\text{Likelihood}} \cdot \underbrace{p(\theta)}_{\text{Prior}}$$

or equivalently,

$$\log p(\theta|x) \propto \sum_{i=1}^N \underbrace{\log f(x_i|\theta)}_{\text{log likelihood}} + \underbrace{\log p(\theta)}_{\text{log prior}}$$

A Stan program specifies the (log) posterior distribution through  
**data**, **parameter** and **model** blocks

## Stan main program blocks

data{...}

parameters{...}

model {...}

- Declaring input data & arguments
- Specifying model parameters to be fitted
- Defining model (log posterior)

$$\log p(\theta|x) \propto \sum_{i=1}^N \underbrace{\log f(x_i|\theta)}_{\text{log likelihood}} + \underbrace{\log p(\theta)}_{\text{log prior}}$$

## Stan program blocks

functions{...}

data{...}

transformed data{...}

parameters{...}

transformed parameters {...}

model{...}

generated quantities{...}

## Example: simple linear regression

Consider the standard simple regression model:

$$Y_i = \alpha + \beta X_i + \epsilon_i, \quad \epsilon_i \stackrel{iid}{\sim} \mathcal{N}(0, \sigma^2)$$

A Bayesian model for this using *weakly informative priors* can be defined as follows

$$Y_i | X_i, \alpha, \beta \sim \mathcal{N}(\alpha + \beta X_i, \sigma^2)$$

$$\alpha \sim \mathcal{N}(0, 10)$$

$$\beta \sim \mathcal{N}(0, 10)$$

$$\sigma \sim \mathcal{N}^+(0, 10)$$

## Example: simple linear regression

```
data {  
    int<lower=0> N;  
    vector[N] x;  
    vector[N] y;  
}  
parameters {  
    real alpha;  
    real beta;  
    real<lower=0> sigma;  
}  
model {  
    y ~ normal(alpha + beta * x, sigma);  
    alpha ~ normal(0, 10);  
    beta ~ normal(0, 10);  
    sigma ~ normal(0, 10);  
}
```

### Comments:

- Variable and types need to be declared
- Each statement ends with ;
- Constraints of variables & parameters enforced using  $\langle \dots \rangle$ -brackets

Most common data types:

- `int`: Integers, eg. discrete data
- `real`: Real numbers, eg. continuous
- `vector`: One dimensional array for storing numbers
- `matrix`: Two dimensional array for storing numbers
- `simplex`: Probability vector (sums to 1)
- Others exist! See manual.
- Array/matrix subsetting same as in standard C/C++

## Example: multiple linear regression

```
data {  
    int<lower=0> N;  
    int<lower=0> K;  
    matrix[N, K] x;  
    vector[N] y;  
}  
parameters {  
    real alpha;  
    vector[K] beta;  
    real<lower=0> sigma;  
}  
model {  
    y ~ normal(x * beta + alpha, sigma);  
    ...  
}
```

## Alternative model definition

Previous example uses vectorized notation for defining likelihood and priors (eg.  $y \sim \text{normal}$ ). A more flexible option is to define the model "term-wise" using a for-loop

```
model {  
  for (i in 1:N){  
    target += lpdf_normal(y[i] | alpha + x[i] * beta, sigma);  
  }  
  target += lpdf_normal(alpha | 0, 10);  
  target += lpdf_normal(beta | 0, 10);  
  ...  
}
```

## Alternative model definition

This alternative notation resembles the mathematical definition of the log posterior

$$\log p(\theta|x) \propto \sum_{i=1}^N \log f(y_i|x_i, \alpha, \beta) + \log p(\alpha) + \log p(\beta) + \log p(\sigma),$$

i.e. `target+=` corresponds adding a log-probability term to the above sum.

## Generated quantities block

Recall, end goal of computing posterior often is to compute

$$\mathbb{E}[h(\theta)] \approx \frac{1}{N} \sum_{i=1}^N h(\theta^{(i)})$$

In Stan, the generated quantities block can be used for implementing the computation of  $h(\theta^{(i)})$  sample-wise. Note: can also be done outside Stan after sampling.

## Generated quantities - example

Example:  $h(\beta) = \beta^2$  and  $h(\beta) = \mathbb{1}(\beta > 0)$

```
generated quantities{
  real beta_sqr =  beta^2;
  int beta_is_significant = beta > 0;
}
```

Most common use of generated quantities is to generate samples from the predictive distribution  $p(\hat{y}|y)$  of new data  $\hat{y}$  given current data  $y$  as will be shown later.

## Model selection

Common likelihood-prior combinations for 1-dim posteriors

Data	Parameter	Likelihood	Prior
Continuous ( $y \in \mathbb{R}$ )	Mean	Normal	Normal
	Variance		Inv-Gamma
Discrete ( $y \in \mathbb{Z}$ )	Mean	Poisson Neg.binomial	Normal
	Variance		Normal
Binary ( $y \in \{0, 1\}$ )	Prob. of success	Bernoulli	Beta
	Treatment effect		Dirichlet Normal

- Prior can both incorporate previous known knowledge (informative) or be chosen to be weakly informative
- Can also be ignored in the Stan program (vague prior)
- See Stan manual for likelihood/prior combinations



- ① Bayesian modeling
- ② Introduction to Stan
- ③ Model diagnostics and evaluation

## Motivation

- How do we know that we can trust that the posterior is computed correctly?
- How do we know if our model/prior is the best possible for our data?
- We will here present common model diagnostic tools for answering these questions

## Sampler convergence - background

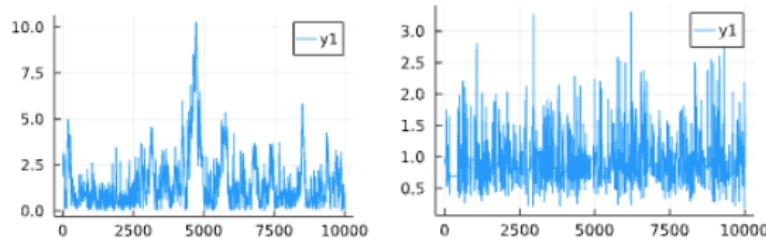
- To be able to "trust" our samples, we need to assure that our sampler has "converged" properly.
- Recall, if we are able produce independent samples  $\theta^{(i)}$  from the true posterior, then as the number of samples  $N$  grows,

$$\frac{1}{N} \sum_{i=1}^N h(\theta^{(i)}) \xrightarrow{a.s} \mathbb{E}[h(\theta)]$$

- Formally: Stan produces  $\theta^{(i)}$  using Markov Chain Monte Carlo (MCMC) methods, by constructing a Markov chain which has the target posterior as its stationary distribution
- Thus, to trust our samples, we must verify that the underlying Markov chain has converged to its stationary distribution

## Traceplots

Idea: verifying convergence by inspecting the samples over time over several chains

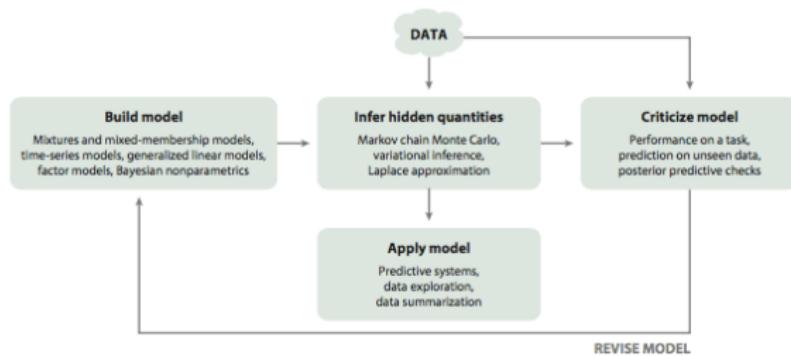


- Left figure: samples clearly correlated, chain not stationary!
- Right: better but not perfect
- Repeat for multiple chains and assure that all chains have found the same range
- Only keep samples after convergence (warm-up phase).

## Quantitative diagnostics - $\hat{R}$ and ESS

- $\hat{R}$ -statistic: comparing variance within and between different chains. Values "close to" 1 indicates that all chains have converged.
- Effective sample size (ESS): The number of independent samples needed for the same uncertainty as in our estimates. Measure of how independent samples are

# Model evaluation - Box's loop



- Formalized by Blei, David M. “Build, compute, critique, repeat: Data analysis with latent variable models.” Annual Review of Statistics and Its Application 1 (2014): 203-232.
- General framework for iterative development of statistical models
- Will now focus on the “Criticize model” block

## Posterior predictive distribution

- Idea: Evaluate/criticise model by how well it predicts new data
- Key object: the posterior predictive distribution (PPD)

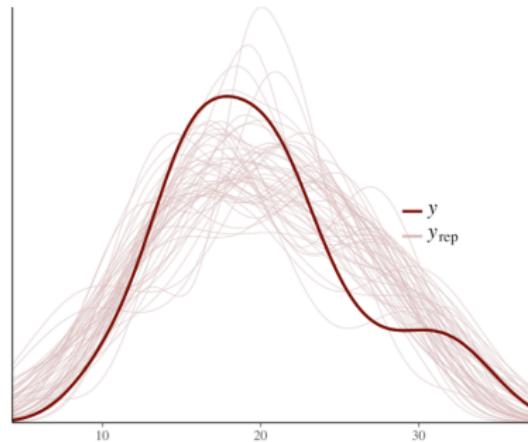
$$p_{post}(\hat{y}|\theta) = \int p(\hat{y}|\theta)p(\theta|y)dy,$$

where  $\hat{y}$  is a new data point

- Posterior predictive checks
- Leave-one-out cross validation

## Posterior predictive check

- Idea: compare samples from the posterior predictive distribution with the actual data



- Useful for identifying obvious mismatches between the model and data
- Can also be done for priors (prior predictive check)

## Sampling from the PPD

- In practice, samples from the PPD can be generated by sampling one data point from the likelihood *for each* posterior sample and data point.
- In Stan, this can be done *in parallel* to the posterior sampling using the generated quantities block as follows

```
generated quantities {
    array[N] real ypred;
    for (i in 1:N){
        ypred[i] = normal_rng(alpha + x[i] * beta, sigma);
    }
}
```

## Sampling from the PPD

- Why does this work?

$$\begin{aligned}\mathbb{E}[h(Y)] &= \int h(\bar{y}) p(\bar{y}|y) dy \\ &= \int \int h(\hat{y}) p(\hat{y}|\theta) p(\theta|y) d\theta dy \\ &\approx \frac{1}{N} \frac{1}{M} \sum_{i=1}^N \sum_{j=1}^M h(\hat{y}^{(i,j)}),\end{aligned}$$

where  $\hat{y}^{(i,j)} \sim p(\hat{y}, \theta^{(j)}|y) = p(\hat{y}|\theta^{(j)})p(\theta^{(j)}|y)$  are samples from the joint distribution corresponding to the sampling procedure from the previous slide.

## Model comparison

- The PPD  $p_{post}$  can also be used for comparing different candidate models
- Useful theoretic metric for this is the expected log predictive density (ELPD)

$$\mathbb{E}[\log(p_{post}(\hat{Y}|\theta))] = \int \log(p_{post}(\hat{y}|\theta))p(\hat{y})d\hat{y},$$

where  $p(\hat{y})$  is the true data generating process

- ELPD measures how well the PPD matches the true distribution on average: if  $p(\hat{y})$  is high,  $\log(p(\hat{y}|\theta))$  should be high and v.v
- However  $p(\hat{y})$  is unknown in general and thus ELPD needs to be estimated → cross validation

## Leave-one-out cross-validation (LOO-CV)

- Leave-one-out cross-validation (LOO-CV) estimates ELPD by

$$\sum_i^n \log p_{post(-i)}(y_i|\theta),$$

where  $p_{post(-i)|\theta}(y_i)$  is the PPD evaluated at  $y_i$  of a posterior computed excluding  $y_i$  as a data point.

- Can be done computationally efficient using Pareto smoothed Importance Sampling (PSIS LOO-CV)
- Implemented in the loo-package

## Computer lab

Now: experiment with Stan on your own computer. Two options:

- A** Try Stan on your own research problem/data
- B** Workshop demo notebook with corresponding exercises at  
<https://github.com/JTorgander/interbayes-workshop>

## References

- [1] Andrew Gelman, John B Carlin, Hal S Stern, and Donald B Rubin.  
*Bayesian data analysis.*  
Chapman and Hall/CRC, 1995.
- [2] Richard McElreath.  
*Statistical rethinking: A Bayesian course with examples in R and Stan.*  
Chapman and Hall/CRC, 2018.